

Dynamic Programming



DEPARTMENT OF COMPUTER SCIENCE
ST. FRANCIS XAVIER UNIVERSITY

CSCI-531 - Reinforcement Learning

Fall 2025

From Theory to Practice: Computing Optimal Policies

What We Know So Far

- ▶ How to model problems as MDPs

From Theory to Practice: Computing Optimal Policies

What We Know So Far

- ▶ How to model problems as MDPs
- ▶ Why we need policies and value functions

From Theory to Practice: Computing Optimal Policies

What We Know So Far

- ▶ How to model problems as MDPs
- ▶ Why we need policies and value functions
- ▶ What optimal policies look like mathematically

From Theory to Practice: Computing Optimal Policies

What We Know So Far

- ▶ How to model problems as MDPs
- ▶ Why we need policies and value functions
- ▶ What optimal policies look like mathematically

The Remaining Challenge

How do we actually compute these optimal policies?

From Theory to Practice: Computing Optimal Policies

What We Know So Far

- ▶ How to model problems as MDPs
- ▶ Why we need policies and value functions
- ▶ What optimal policies look like mathematically

The Remaining Challenge

How do we actually compute these optimal policies?

The Approach Menu

- ▶ Dynamic Programming: Complete knowledge of MDP

From Theory to Practice: Computing Optimal Policies

What We Know So Far

- ▶ How to model problems as MDPs
- ▶ Why we need policies and value functions
- ▶ What optimal policies look like mathematically

The Remaining Challenge

How do we actually compute these optimal policies?

The Approach Menu

- ▶ Dynamic Programming: Complete knowledge of MDP
- ▶ Monte Carlo Methods: Learn from experience

From Theory to Practice: Computing Optimal Policies

What We Know So Far

- ▶ How to model problems as MDPs
- ▶ Why we need policies and value functions
- ▶ What optimal policies look like mathematically

The Remaining Challenge

How do we actually compute these optimal policies?

The Approach Menu

- ▶ Dynamic Programming: Complete knowledge of MDP
- ▶ Monte Carlo Methods: Learn from experience
- ▶ Temporal Difference Learning: Combines both approaches

Why Start with Dynamic Programming?

Three Key Reasons

1. Conceptual Foundation

- ▶ Core principles used in all RL algorithms

Why Start with Dynamic Programming?

Three Key Reasons

1. Conceptual Foundation

- ▶ Core principles used in all RL algorithms

2. Theoretical Clarity

- ▶ Shows exactly what we're trying to approximate

Why Start with Dynamic Programming?

Three Key Reasons

1. Conceptual Foundation

- ▶ Core principles used in all RL algorithms

2. Theoretical Clarity

- ▶ Shows exactly what we're trying to approximate

3. Practical Utility

- ▶ Works perfectly when MDP is known

Why Start with Dynamic Programming?

Three Key Reasons

1. Conceptual Foundation


- ▶ Core principles used in all RL algorithms

2. Theoretical Clarity

- ▶ Shows exactly what we're trying to approximate

3. Practical Utility

- ▶ Works perfectly when MDP is known



Why Start with Dynamic Programming?

Three Key Reasons

1. Conceptual Foundation


- ▶ Core principles used in all RL algorithms

2. Theoretical Clarity

- ▶ Shows exactly what we're trying to approximate

3. Practical Utility

- ▶ Works perfectly when MDP is known



What do we need to calculate to obtain optimal policies?

Policy Iteration: The Intuitive Approach

Core Idea

Start with any policy, then repeatedly improve it until optimal

Policy Iteration: The Intuitive Approach

Core Idea

Start with any policy, then repeatedly improve it until optimal

Two Alternating Steps

1. **Policy Evaluation:** How good is current policy?

Policy Iteration: The Intuitive Approach

Core Idea

Start with any policy, then repeatedly improve it until optimal

Two Alternating Steps

1. **Policy Evaluation**: How good is current policy?
2. **Policy Improvement**: Update policy based on values


Policy Iteration: The Intuitive Approach

Core Idea

Start with any policy, then repeatedly improve it until optimal

Two Alternating Steps

1. **Policy Evaluation:** How good is current policy?
2. **Policy Improvement:** Update policy based on values




Policy Iteration: The Intuitive Approach

Core Idea

Start with any policy, then repeatedly improve it until optimal

Two Alternating Steps

1. **Policy Evaluation**: How good is current policy?
2. **Policy Improvement**: Update policy based on values



Each improvement step makes the policy better (or keeps it optimal)

Policy Evaluation: Computing Value Functions

The Challenge

Given policy π , compute $v_\pi(s)$ for all states using:

$$v_\pi(s) = \sum_a \pi(a|s) \sum_{s'} p(s'|s, a) [r(s, a) + \gamma v_\pi(s')]$$

Policy Evaluation: Computing Value Functions

The Challenge

Given policy π , compute $v_\pi(s)$ for all states using:

$$v_\pi(s) = \sum_a \pi(a|s) \sum_{s'} p(s'|s, a) [r(s, a) + \gamma v_\pi(s')]$$

The Problem

- ▶ This equation is **circular!**

Policy Evaluation: Computing Value Functions

The Challenge

Given policy π , compute $v_\pi(s)$ for all states using:

$$v_\pi(s) = \sum_a \pi(a|s) \sum_{s'} p(s'|s, a) [r(s, a) + \gamma v_\pi(s')]$$

The Problem

- ▶ This equation is **circular**!
- ▶ Value of state s depends on values of future states s'

Policy Evaluation: Computing Value Functions

The Challenge

Given policy π , compute $v_\pi(s)$ for all states using:

$$v_\pi(s) = \sum_a \pi(a|s) \sum_{s'} p(s'|s, a) [r(s, a) + \gamma v_\pi(s')]$$

The Problem

- ▶ This equation is **circular**!
- ▶ Value of state s depends on values of future states s'
- ▶ We have a system of $|S|$ equations with $|S|$ unknowns

Policy Evaluation: Computing Value Functions

The Challenge

Given policy π , compute $v_\pi(s)$ for all states using:

$$v_\pi(s) = \sum_a \pi(a|s) \sum_{s'} p(s'|s, a) [r(s, a) + \gamma v_\pi(s')]$$

Policy Evaluation: Computing Value Functions

The Challenge

Given policy π , compute $v_\pi(s)$ for all states using:

$$v_\pi(s) = \sum_a \pi(a|s) \sum_{s'} p(s'|s, a) [r(s, a) + \gamma v_\pi(s')]$$

The Solution: Iterative Approximation

1. Start with initial guesses $v_0(s)$ (often zeros)

Policy Evaluation: Computing Value Functions

The Challenge

Given policy π , compute $v_\pi(s)$ for all states using:

$$v_\pi(s) = \sum_a \pi(a|s) \sum_{s'} p(s'|s, a) [r(s, a) + \gamma v_\pi(s')]$$

The Solution: Iterative Approximation

1. Start with initial guesses $v_0(s)$ (often zeros)
2. Iteratively improve estimates using Bellman equation

Policy Evaluation: Computing Value Functions

The Challenge

Given policy π , compute $v_\pi(s)$ for all states using:

$$v_\pi(s) = \sum_a \pi(a|s) \sum_{s'} p(s'|s, a) [r(s, a) + \gamma v_\pi(s')]$$

The Solution: Iterative Approximation


1. Start with initial guesses $v_0(s)$ (often zeros)
2. Iteratively improve estimates using Bellman equation
3. Continue until values stabilize

Policy Evaluation Algorithm

Algorithm Policy Evaluation


- 1: **Initialize:** $V(s) \in \mathbb{R}$ arbitrarily, $V(\text{terminal}) = 0$
- 2: **repeat**
- 3: $\Delta \leftarrow 0$
- 4: **for** each state s **do**
- 5: $v \leftarrow V(s)$
- 6: $V(s) \leftarrow \sum_a \pi(a|s) \sum_{s'} p(s'|s, a)[r + \gamma V(s')]$
- 7: $\Delta \leftarrow \max(\Delta, |v - V(s)|)$
- 8: **end for**
- 9: **until** $\Delta < \theta$ (small threshold)

Policy Evaluation: Intuition



Iteration 0: Initialize all values to 0


Policy Evaluation: Intuition



Iteration 0: Initialize all values to 0

Iteration 1: Values propagate to immediate neighbors

Policy Evaluation: Intuition




Iteration 0: Initialize all values to 0

Iteration 1: Values propagate to immediate neighbors

Iteration 2: Values continue propagating (discounted)

Policy Evaluation: Intuition



Iteration 0: Initialize all values to 0

Iteration 1: Values propagate to immediate neighbors

Iteration 2: Values continue propagating (discounted)

Convergence: Optimal values for given policy

Policy Improvement: Making Better Decisions

The Question

Given policy π and its value function v_π , can we find a better policy?

Policy Improvement: Making Better Decisions

The Question

Given policy π and its value function v_π , can we find a better policy?

The Approach: Be Greedy!

For each state s , examine all possible actions:

$$q_\pi(s, a) = \sum_{s'} p(s'|s, a) [r(s, a) + \gamma v_\pi(s')]$$

"What happens if I take action a in state s , then follow policy π ?"

Policy Improvement: Making Better Decisions

The Question

Given policy π and its value function v_π , can we find a better policy?

The Approach: Be Greedy!

For each state s , examine all possible actions:

$$q_\pi(s, a) = \sum_{s'} p(s'|s, a) [r(s, a) + \gamma v_\pi(s')]$$

"What happens if I take action a in state s , then follow policy π ?"

Improved Policy

Create new policy π' by being greedy:

$$\pi'(s) = \arg \max_a q_\pi(s, a) = \arg \max_a \sum_{s'} p(s'|s, a) [r(s, a) + \gamma v_\pi(s')]$$

Policy Improvement Theorem

Theorem

If π and π' are policies such that for all $s \in S$:

$$q_{\pi}(s, \pi'(s)) \geq v_{\pi}(s)$$

Then policy π' must be as good as or better than π :

$$v_{\pi'}(s) \geq v_{\pi}(s) \text{ for all } s \in S$$

Policy Improvement Theorem

Theorem

If π and π' are policies such that for all $s \in S$:

$$q_{\pi}(s, \pi'(s)) \geq v_{\pi}(s)$$

Then policy π' must be as good as or better than π :

$$v_{\pi'}(s) \geq v_{\pi}(s) \text{ for all } s \in S$$

Intuition

- ▶ Acting optimally for one step, then following old policy

Policy Improvement Theorem

Theorem

If π and π' are policies such that for all $s \in S$:

$$q_{\pi}(s, \pi'(s)) \geq v_{\pi}(s)$$

Then policy π' must be as good as or better than π :

$$v_{\pi'}(s) \geq v_{\pi}(s) \text{ for all } s \in S$$

Intuition

- ▶ Acting optimally for one step, then following old policy
- ▶ Can only make things **better** (or stay the same)

Policy Improvement Theorem

Theorem

If π and π' are policies such that for all $s \in S$:

$$q_{\pi}(s, \pi'(s)) \geq v_{\pi}(s)$$

Then policy π' must be as good as or better than π :

$$v_{\pi'}(s) \geq v_{\pi}(s) \text{ for all } s \in S$$

Intuition

- ▶ Acting optimally for one step, then following old policy
- ▶ Can only make things better (or stay the same)

Guaranteed improvement or stability!

Policy Iteration Algorithm

Complete Process

1. **Initialize:** Start with arbitrary policy π_0

Policy Iteration Algorithm

Complete Process

1. **Initialize:** Start with arbitrary policy π_0
2. **Policy Evaluation:** Compute v_π for current policy

Policy Iteration Algorithm

Complete Process

1. **Initialize:** Start with arbitrary policy π_0
2. **Policy Evaluation:** Compute v_π for current policy
3. **Policy Improvement:** Compute improved policy π' from v_π

Policy Iteration Algorithm

Complete Process

1. **Initialize:** Start with arbitrary policy π_0
2. **Policy Evaluation:** Compute v_π for current policy
3. **Policy Improvement:** Compute improved policy π' from v_π
4. **Check:** If $\pi' = \pi$, STOP (found π^*)

Policy Iteration Algorithm

Complete Process

1. **Initialize:** Start with arbitrary policy π_0
2. **Policy Evaluation:** Compute v_π for current policy
3. **Policy Improvement:** Compute improved policy π' from v_π
4. **Check:** If $\pi' = \pi$, STOP (found π^*)
5. **Otherwise:** Set $\pi = \pi'$ and go to step 2

Policy Iteration Algorithm

Complete Process

1. **Initialize:** Start with arbitrary policy π_0
2. **Policy Evaluation:** Compute v_π for current policy
3. **Policy Improvement:** Compute improved policy π' from v_π
4. **Check:** If $\pi' = \pi$, STOP (found π^*)
5. **Otherwise:** Set $\pi = \pi'$ and go to step 2

Convergence Guarantee

- ▶ MDPs have finite number of deterministic policies

Policy Iteration Algorithm

Complete Process

1. **Initialize:** Start with arbitrary policy π_0
2. **Policy Evaluation:** Compute v_π for current policy
3. **Policy Improvement:** Compute improved policy π' from v_π
4. **Check:** If $\pi' = \pi$, STOP (found π^*)
5. **Otherwise:** Set $\pi = \pi'$ and go to step 2

Convergence Guarantee

- ▶ MDPs have finite number of deterministic policies
- ▶ Each improvement makes policy strictly better (unless optimal)

Policy Iteration Algorithm

Complete Process


1. **Initialize:** Start with arbitrary policy π_0
2. **Policy Evaluation:** Compute v_π for current policy
3. **Policy Improvement:** Compute improved policy π' from v_π
4. **Check:** If $\pi' = \pi$, STOP (found π^*)
5. **Otherwise:** Set $\pi = \pi'$ and go to step 2

Convergence Guarantee

- ▶ MDPs have finite number of deterministic policies
- ▶ Each improvement makes policy strictly better (unless optimal)
- ▶ **Must converge to optimal policy in finite steps**

Policy Iteration Example


Step 1: Random Policy



Random actions

Policy Iteration Example

Step 1: Random Policy



Random actions


Step 2: Evaluate



Computed values

Policy Iteration Example

Step 1: Random Policy




Random actions

Step 2: Evaluate



Computed values


Step 3: Improve



Greedy actions

Policy Iteration Example

Step 1: Random Policy




Random actions

Step 2: Evaluate



Computed values

Step 3: Improve



Greedy actions

Key Insight

Each iteration makes the policy strictly better until optimal

Value Iteration: Cutting the Bottleneck

Policy Iteration's Bottleneck

- ▶ Full policy evaluation requires multiple sweeps through all states

Value Iteration: Cutting the Bottleneck

Policy Iteration's Bottleneck

- ▶ Full policy evaluation requires multiple sweeps through all states
- ▶ Each evaluation step can be **computationally expensive**

Value Iteration: Cutting the Bottleneck

Policy Iteration's Bottleneck

- ▶ Full policy evaluation requires multiple sweeps through all states
- ▶ Each evaluation step can be **computationally expensive**
- ▶ We do complete evaluation before any improvement

Value Iteration: Cutting the Bottleneck

Policy Iteration's Bottleneck

- ▶ Full policy evaluation requires multiple sweeps through all states
- ▶ Each evaluation step can be **computationally expensive**
- ▶ We do complete evaluation before any improvement

Value Iteration's Insight

What if we combine evaluation and improvement?

Value Iteration: Cutting the Bottleneck

Policy Iteration's Bottleneck

- ▶ Full policy evaluation requires multiple sweeps through all states
- ▶ Each evaluation step can be **computationally expensive**
- ▶ We do complete evaluation before any improvement

Value Iteration's Insight

What if we combine evaluation and improvement?

- ▶ Don't need full evaluation before improving

Value Iteration: Cutting the Bottleneck

Policy Iteration's Bottleneck

- ▶ Full policy evaluation requires multiple sweeps through all states
- ▶ Each evaluation step can be **computationally expensive**
- ▶ We do complete evaluation before any improvement

Value Iteration's Insight

What if we combine evaluation and improvement?

- ▶ Don't need full evaluation before improving
- ▶ One step of evaluation gives useful information

Value Iteration: Cutting the Bottleneck

Policy Iteration's Bottleneck

- ▶ Full policy evaluation requires multiple sweeps through all states
- ▶ Each evaluation step can be **computationally expensive**
- ▶ We do complete evaluation before any improvement

Value Iteration's Insight

What if we combine evaluation and improvement?

- ▶ Don't need full evaluation before improving
- ▶ One step of evaluation gives useful information
- ▶ Always act greedily based on current value estimates

Value Iteration: Cutting the Bottleneck

Policy Iteration's Bottleneck

- ▶ Full policy evaluation requires multiple sweeps through all states
- ▶ Each evaluation step can be **computationally expensive**
- ▶ We do complete evaluation before any improvement

Value Iteration's Insight

What if we combine evaluation and improvement?

- ▶ Don't need full evaluation before improving
- ▶ One step of evaluation gives useful information
- ▶ Always act greedily based on current value estimates

Skip the separate policy - work directly with values!

Value Iteration Formula

Key Difference

Instead of following a fixed policy, always choose the best action:

$$v_{k+1}(s) = \max_a \sum_{s'} p(s'|s, a) [r(s, a) + \gamma v_k(s')]$$

Value Iteration Formula

Key Difference

Instead of following a fixed policy, always choose the best action:

$$v_{k+1}(s) = \max_a \sum_{s'} p(s'|s, a) [r(s, a) + \gamma v_k(s')]$$

Comparison

Policy Iteration says:

- ▶ "Fully evaluate my current strategy, then improve it"

Value Iteration says:

Value Iteration Formula

Key Difference

Instead of following a fixed policy, always choose the best action:

$$v_{k+1}(s) = \max_a \sum_{s'} p(s'|s, a) [r(s, a) + \gamma v_k(s')]$$

Comparison

Policy Iteration says:

- ▶ "Fully evaluate my current strategy, then improve it"

Value Iteration says:

- ▶ "Always act greedily based on current value estimates"

Value Iteration Formula

Key Difference

Instead of following a fixed policy, always choose the best action:

$$v_{k+1}(s) = \max_a \sum_{s'} p(s'|s, a) [r(s, a) + \gamma v_k(s')]$$

Value Iteration Formula

Key Difference

Instead of following a fixed policy, always choose the best action:

$$v_{k+1}(s) = \max_a \sum_{s'} p(s'|s, a) [r(s, a) + \gamma v_k(s')]$$

Connection to Bellman Optimality

This directly solves the Bellman optimality equation:

$$v^*(s) = \max_a \sum_{s'} p(s'|s, a) [r(s, a) + \gamma v^*(s')]$$

Value Iteration Algorithm

Value Iteration Algorithm

1. **Initialize:** $V(s) \in \mathbb{R}$ arbitrarily for all $s \in S$, $V(\text{terminal}) = 0$

Value Iteration Algorithm

Value Iteration Algorithm

1. **Initialize:** $V(s) \in \mathbb{R}$ arbitrarily for all $s \in S$, $V(\text{terminal}) = 0$
2. **Repeat** until convergence:
 - ▶ $\Delta \leftarrow 0$

Value Iteration Algorithm

Value Iteration Algorithm

1. **Initialize:** $V(s) \in \mathbb{R}$ arbitrarily for all $s \in S$, $V(\text{terminal}) = 0$
2. **Repeat** until convergence:
 - ▶ $\Delta \leftarrow 0$
 - ▶ **For** each state $s \in S$:
 - ▶ $v \leftarrow V(s)$

Value Iteration Algorithm

Value Iteration Algorithm

1. **Initialize:** $V(s) \in \mathbb{R}$ arbitrarily for all $s \in S$, $V(\text{terminal}) = 0$
2. **Repeat** until convergence:
 - ▶ $\Delta \leftarrow 0$
 - ▶ **For** each state $s \in S$:
 - ▶ $v \leftarrow V(s)$
 - ▶ $V(s) \leftarrow \max_a \sum_{s'} p(s'|s, a)[r + \gamma V(s')]$

Value Iteration Algorithm

Value Iteration Algorithm

1. **Initialize:** $V(s) \in \mathbb{R}$ arbitrarily for all $s \in S$, $V(\text{terminal}) = 0$
2. **Repeat** until convergence:
 - ▶ $\Delta \leftarrow 0$
 - ▶ **For** each state $s \in S$:
 - ▶ $v \leftarrow V(s)$
 - ▶ $V(s) \leftarrow \max_a \sum_{s'} p(s'|s, a)[r + \gamma V(s')]$
 - ▶ $\Delta \leftarrow \max(\Delta, |v - V(s)|)$

Value Iteration Algorithm

Value Iteration Algorithm

1. **Initialize:** $V(s) \in \mathbb{R}$ arbitrarily for all $s \in S$, $V(\text{terminal}) = 0$
2. **Repeat** until convergence:
 - ▶ $\Delta \leftarrow 0$
 - ▶ **For** each state $s \in S$:
 - ▶ $v \leftarrow V(s)$
 - ▶ $V(s) \leftarrow \max_a \sum_{s'} p(s'|s, a)[r + \gamma V(s')]$
 - ▶ $\Delta \leftarrow \max(\Delta, |v - V(s)|)$
3. **Until** $\Delta < \theta$ (convergence threshold)

Value Iteration Algorithm

Value Iteration Algorithm

1. **Initialize:** $V(s) \in \mathbb{R}$ arbitrarily for all $s \in S$, $V(\text{terminal}) = 0$
2. **Repeat** until convergence:
 - ▶ $\Delta \leftarrow 0$
 - ▶ **For** each state $s \in S$:
 - ▶ $v \leftarrow V(s)$
 - ▶ $V(s) \leftarrow \max_a \sum_{s'} p(s'|s, a)[r + \gamma V(s')]$
 - ▶ $\Delta \leftarrow \max(\Delta, |v - V(s)|)$
3. **Until** $\Delta < \theta$ (convergence threshold)
4. **Extract policy:** $\pi(s) = \arg \max_a \sum_{s'} p(s'|s, a)[r + \gamma V(s')]$

Value Iteration Algorithm

Value Iteration Algorithm

1. **Initialize:** $V(s) \in \mathbb{R}$ arbitrarily for all $s \in S$, $V(\text{terminal}) = 0$
2. **Repeat** until convergence:
 - ▶ $\Delta \leftarrow 0$
 - ▶ **For** each state $s \in S$:
 - ▶ $v \leftarrow V(s)$
 - ▶ $V(s) \leftarrow \max_a \sum_{s'} p(s'|s, a)[r + \gamma V(s')]$
 - ▶ $\Delta \leftarrow \max(\Delta, |v - V(s)|)$
3. **Until** $\Delta < \theta$ (convergence threshold)
4. **Extract policy:** $\pi(s) = \arg \max_a \sum_{s'} p(s'|s, a)[r + \gamma V(s')]$

Key Advantages

- ▶ **Simpler:** No separate policy evaluation phase

Value Iteration Algorithm

Value Iteration Algorithm

1. **Initialize:** $V(s) \in \mathbb{R}$ arbitrarily for all $s \in S$, $V(\text{terminal}) = 0$
2. **Repeat** until convergence:
 - ▶ $\Delta \leftarrow 0$
 - ▶ **For** each state $s \in S$:
 - ▶ $v \leftarrow V(s)$
 - ▶ $V(s) \leftarrow \max_a \sum_{s'} p(s'|s, a)[r + \gamma V(s')]$
 - ▶ $\Delta \leftarrow \max(\Delta, |v - V(s)|)$
3. **Until** $\Delta < \theta$ (convergence threshold)
4. **Extract policy:** $\pi(s) = \arg \max_a \sum_{s'} p(s'|s, a)[r + \gamma V(s')]$

Key Advantages

- ▶ **Simpler:** No separate policy evaluation phase
- ▶ **More efficient:** Often faster per iteration

Value Iteration vs Policy Iteration

Aspect	Policy Iteration	Value Iteration
Approach	Evaluate then improve	Combined eval+improve
Per iteration	Slower (full eval)	Faster (one step)
Total iterations	Fewer	More
Memory	Store policy + values	Store values only
Implementation	More complex	Simpler
Convergence	Finite steps	Asymptotic

Value Iteration vs Policy Iteration

Aspect	Policy Iteration	Value Iteration
Approach	Evaluate then improve	Combined eval+improve
Per iteration	Slower (full eval)	Faster (one step)
Total iterations	Fewer	More
Memory	Store policy + values	Store values only
Implementation	More complex	Simpler
Convergence	Finite steps	Asymptotic

When to Use Which?

- ▶ **Policy Iteration:** When actions are expensive to compute

Value Iteration vs Policy Iteration

Aspect	Policy Iteration	Value Iteration
Approach	Evaluate then improve	Combined eval+improve
Per iteration	Slower (full eval)	Faster (one step)
Total iterations	Fewer	More
Memory	Store policy + values	Store values only
Implementation	More complex	Simpler
Convergence	Finite steps	Asymptotic

When to Use Which?

- ▶ **Policy Iteration:** When actions are expensive to compute
- ▶ **Value Iteration:** When state space is large

Value Iteration vs Policy Iteration

Aspect	Policy Iteration	Value Iteration
Approach	Evaluate then improve	Combined eval+improve
Per iteration	Slower (full eval)	Faster (one step)
Total iterations	Fewer	More
Memory	Store policy + values	Store values only
Implementation	More complex	Simpler
Convergence	Finite steps	Asymptotic

When to Use Which?

- ▶ **Policy Iteration:** When actions are expensive to compute
- ▶ **Value Iteration:** When state space is large
- ▶ **In practice:** Value iteration often preferred for simplicity

Computational Complexity Analysis

Time Complexity per Iteration

For both algorithms, each iteration requires:

- ▶ Loop over all states: $O(|S|)$

Computational Complexity Analysis

Time Complexity per Iteration

For both algorithms, each iteration requires:

- ▶ Loop over all states: $O(|S|)$
- ▶ For each state, loop over all actions: $O(|A|)$

Computational Complexity Analysis

Time Complexity per Iteration

For both algorithms, each iteration requires:

- ▶ Loop over all states: $O(|S|)$
- ▶ For each state, loop over all actions: $O(|A|)$
- ▶ For each action, loop over possible next states: $O(|S|)$

Computational Complexity Analysis

Time Complexity per Iteration

For both algorithms, each iteration requires:

- ▶ Loop over all states: $O(|S|)$
- ▶ For each state, loop over all actions: $O(|A|)$
- ▶ For each action, loop over possible next states: $O(|S|)$
- ▶ **Total per iteration:** $O(|S|^2|A|)$

Computational Complexity Analysis

Time Complexity per Iteration

For both algorithms, each iteration requires:

- ▶ Loop over all states: $O(|S|)$
- ▶ For each state, loop over all actions: $O(|A|)$
- ▶ For each action, loop over possible next states: $O(|S|)$
- ▶ **Total per iteration:** $O(|S|^2|A|)$

Space Complexity

- ▶ **Value Iteration:** $O(|S|)$ for value function

Computational Complexity Analysis

Time Complexity per Iteration

For both algorithms, each iteration requires:

- ▶ Loop over all states: $O(|S|)$
- ▶ For each state, loop over all actions: $O(|A|)$
- ▶ For each action, loop over possible next states: $O(|S|)$
- ▶ **Total per iteration:** $O(|S|^2|A|)$

Space Complexity

- ▶ **Value Iteration:** $O(|S|)$ for value function
- ▶ **Policy Iteration:** $O(|S|)$ for values + $O(|S|)$ for policy

Computational Complexity Analysis

Time Complexity per Iteration

For both algorithms, each iteration requires:

- ▶ Loop over all states: $O(|S|)$
- ▶ For each state, loop over all actions: $O(|A|)$
- ▶ For each action, loop over possible next states: $O(|S|)$
- ▶ **Total per iteration:** $O(|S|^2|A|)$

Space Complexity

- ▶ **Value Iteration:** $O(|S|)$ for value function
- ▶ **Policy Iteration:** $O(|S|)$ for values + $O(|S|)$ for policy
- ▶ Both need to store MDP: $O(|S|^2|A|)$ for transition probabilities

Computational Complexity Analysis

Scalability Issues

- ▶ **Curse of dimensionality:** State space grows exponentially
- ▶ **Memory requirements:** Storing full transition model
- ▶ **Exact methods:** Limited to small-medium problems

Handling Large State Spaces

The Problem

- ▶ Tabular DP requires storing $V(s)$ for every state

Handling Large State Spaces

The Problem

- ▶ Tabular DP requires storing $V(s)$ for every state
- ▶ Real problems often have millions/billions of states

Handling Large State Spaces

The Problem

- ▶ Tabular DP requires storing $V(s)$ for every state
- ▶ Real problems often have millions/billions of states
- ▶ **Solution:** Approximate value function

Handling Large State Spaces

The Problem

- ▶ Tabular DP requires storing $V(s)$ for every state
- ▶ Real problems often have millions/billions of states
- ▶ **Solution:** Approximate value function

Function Approximation Approaches

- ▶ **Linear:** $V(s) \approx \theta^T \phi(s)$
 - ▶ Feature vector $\phi(s)$, learned weights θ

Handling Large State Spaces

The Problem

- ▶ Tabular DP requires storing $V(s)$ for every state
- ▶ Real problems often have millions/billions of states
- ▶ **Solution:** Approximate value function

Function Approximation Approaches

- ▶ **Linear:** $V(s) \approx \theta^T \phi(s)$
 - ▶ Feature vector $\phi(s)$, learned weights θ
- ▶ **Neural Networks:** $V(s) \approx f_\theta(s)$
 - ▶ Deep networks for complex state representations

Handling Large State Spaces

The Problem

- ▶ Tabular DP requires storing $V(s)$ for every state
- ▶ Real problems often have millions/billions of states
- ▶ **Solution:** Approximate value function

Function Approximation Approaches

- ▶ **Linear:** $V(s) \approx \theta^T \phi(s)$
 - ▶ Feature vector $\phi(s)$, learned weights θ
- ▶ **Neural Networks:** $V(s) \approx f_\theta(s)$
 - ▶ Deep networks for complex state representations
- ▶ **Decision Trees:** Partition state space

From Dynamic Programming to Model-Free RL

DP Limitations in Practice

- ▶ **Model requirement:** Often don't know T and R

From Dynamic Programming to Model-Free RL

DP Limitations in Practice

- ▶ **Model requirement:** Often don't know T and R
- ▶ **Computational cost:** Exponential in state space

From Dynamic Programming to Model-Free RL

DP Limitations in Practice

- ▶ **Model requirement:** Often don't know T and R
- ▶ **Computational cost:** Exponential in state space

From Dynamic Programming to Model-Free RL

DP Limitations in Practice

- ▶ **Model requirement:** Often don't know T and R
- ▶ **Computational cost:** Exponential in state space

The Next Step: Model-Free Methods

- ▶ **Monte Carlo:** Learn from complete episodes

From Dynamic Programming to Model-Free RL

DP Limitations in Practice

- ▶ **Model requirement:** Often don't know T and R
- ▶ **Computational cost:** Exponential in state space

The Next Step: Model-Free Methods

- ▶ **Monte Carlo:** Learn from complete episodes
- ▶ **Temporal Difference:** Learn from single steps

From Dynamic Programming to Model-Free RL

DP Limitations in Practice

- ▶ **Model requirement:** Often don't know T and R
- ▶ **Computational cost:** Exponential in state space

The Next Step: Model-Free Methods

- ▶ **Monte Carlo:** Learn from complete episodes
- ▶ **Temporal Difference:** Learn from single steps
- ▶ **Q-Learning:** Off-policy value learning

From Dynamic Programming to Model-Free RL

DP Limitations in Practice

- ▶ **Model requirement:** Often don't know T and R
- ▶ **Computational cost:** Exponential in state space

The Next Step: Model-Free Methods

- ▶ **Monte Carlo:** Learn from complete episodes
- ▶ **Temporal Difference:** Learn from single steps
- ▶ **Q-Learning:** Off-policy value learning
- ▶ **SARSA:** On-policy value learning

From Dynamic Programming to Model-Free RL

DP Limitations in Practice

- ▶ **Model requirement:** Often don't know T and R
- ▶ **Computational cost:** Exponential in state space

The Next Step: Model-Free Methods

- ▶ **Monte Carlo:** Learn from complete episodes
- ▶ **Temporal Difference:** Learn from single steps
- ▶ **Q-Learning:** Off-policy value learning
- ▶ **SARSA:** On-policy value learning

Key Insight

These methods use the same core principles as DP, but learn from experience instead of using a model

Summary

What We've Learned

- ▶ **Policy Iteration:** Alternate evaluation and improvement
- ▶ **Value Iteration:** Combined evaluation and improvement
- ▶ **Theoretical guarantees:** Convergence to optimal policies
- ▶ **Practical considerations:** Complexity and scalability

Next Topics

- ▶ Monte Carlo methods for model-free learning
- ▶ Temporal difference learning (TD, Q-Learning, SARSA)
- ▶ Function approximation and deep reinforcement learning
- ▶ Policy gradient methods