

Approximation - On-Policy Prediction



DEPARTMENT OF COMPUTER SCIENCE
ST. FRANCIS XAVIER UNIVERSITY

CSCI-531 - Reinforcement Learning

Fall 2025

The Tabular Method Approach

What are Tabular Methods?

- Store **exact value** $q_{\pi}(s, a)$ for **each** state-action pair

Q-Table

(s_1, a_1)	$q(s_1, a_1)$
(s_1, a_2)	$q(s_1, a_2)$
\vdots	\vdots
(s_n, a_m)	$q(s_n, a_m)$

The Tabular Method Approach

What are Tabular Methods?

- ▶ Store **exact value** $q_{\pi}(s, a)$ for **each** state-action pair
- ▶ Create a **lookup table** with all possible combinations

Q-Table

(s_1, a_1)	$q(s_1, a_1)$
(s_1, a_2)	$q(s_1, a_2)$
\vdots	\vdots
(s_n, a_m)	$q(s_n, a_m)$

The Tabular Method Approach

What are Tabular Methods?

- ▶ Store **exact value** $q_{\pi}(s, a)$ for **each** state-action pair
- ▶ Create a **lookup table** with all possible combinations
- ▶ Perfect representation when feasible

Q-Table

(s_1, a_1)	$q(s_1, a_1)$
(s_1, a_2)	$q(s_1, a_2)$
\vdots	\vdots
(s_n, a_m)	$q(s_n, a_m)$

The Tabular Method Approach

What are Tabular Methods?

- ▶ Store **exact value** $q_{\pi}(s, a)$ for **each** state-action pair
- ▶ Create a **lookup table** with all possible combinations
- ▶ Perfect representation when feasible

Q-Table

(s_1, a_1)	$q(s_1, a_1)$
(s_1, a_2)	$q(s_1, a_2)$
\vdots	\vdots
(s_n, a_m)	$q(s_n, a_m)$

Activity

What issues can arise with this approach for large state spaces?

Problems with Tabular Methods

Scalability Issues

- ▶ **Memory explosion:** Space complexity $O(|S| \times |A|)$

Problems with Tabular Methods

Scalability Issues

- ▶ **Memory explosion:** Space complexity $O(|S| \times |A|)$
- ▶ **Learning time:** Need to visit every state-action pair

Problems with Tabular Methods

Scalability Issues

- ▶ **Memory explosion:** Space complexity $O(|S| \times |A|)$
- ▶ **Learning time:** Need to visit every state-action pair
- ▶ **Generalization:** No knowledge transfer between similar states

Problems with Tabular Methods

Scalability Issues

- ▶ **Memory explosion:** Space complexity $O(|S| \times |A|)$
- ▶ **Learning time:** Need to visit every state-action pair
- ▶ **Generalization:** No knowledge transfer between similar states

Real-World Examples

- ▶ **Chess:** $\approx 10^{47}$ possible positions

Problems with Tabular Methods

Scalability Issues

- ▶ **Memory explosion:** Space complexity $O(|S| \times |A|)$
- ▶ **Learning time:** Need to visit every state-action pair
- ▶ **Generalization:** No knowledge transfer between similar states

Real-World Examples

- ▶ **Chess:** $\approx 10^{47}$ possible positions
- ▶ **Atari games:** $256^{84 \times 84} \approx 10^{170,000}$ pixel combinations

Problems with Tabular Methods

Scalability Issues

- ▶ **Memory explosion:** Space complexity $O(|S| \times |A|)$
- ▶ **Learning time:** Need to visit every state-action pair
- ▶ **Generalization:** No knowledge transfer between similar states

Real-World Examples

- ▶ **Chess:** $\approx 10^{47}$ possible positions
- ▶ **Atari games:** $256^{84 \times 84} \approx 10^{170,000}$ pixel combinations
- ▶ **Autonomous driving:** Continuous sensor readings

Problems with Tabular Methods

Scalability Issues

- ▶ **Memory explosion:** Space complexity $O(|S| \times |A|)$
- ▶ **Learning time:** Need to visit every state-action pair
- ▶ **Generalization:** No knowledge transfer between similar states

Real-World Examples


- ▶ **Chess:** $\approx 10^{47}$ possible positions
- ▶ **Atari games:** $256^{84 \times 84} \approx 10^{170,000}$ pixel combinations
- ▶ **Autonomous driving:** Continuous sensor readings

Solution: Function Approximation

The Fundamental Shift: Tables vs Functions

Tabular Method (What We've Done Before)

- **One number per state:** Store $v(s_1), v(s_2), \dots, v(s_n)$




Tabular: Each state has its own value

The Fundamental Shift: Tables vs Functions

Tabular Method (What We've Done Before)

- ▶ **One number per state:** Store $v(s_1), v(s_2), \dots, v(s_n)$
- ▶ **Independent values:** Updating $v(s_1)$ doesn't affect $v(s_2)$




Tabular: Each state has its own value

The Fundamental Shift: Tables vs Functions

Tabular Method (What We've Done Before)

- ▶ **One number per state:** Store $v(s_1), v(s_2), \dots, v(s_n)$
- ▶ **Independent values:** Updating $v(s_1)$ doesn't affect $v(s_2)$
- ▶ **Perfect but expensive:** Exact values but need $|S|$ memory slots




Tabular: Each state has its own value

The Fundamental Shift: Tables vs Functions

Tabular Method (What We've Done Before)

- ▶ **One number per state:** Store $v(s_1), v(s_2), \dots, v(s_n)$
- ▶ **Independent values:** Updating $v(s_1)$ doesn't affect $v(s_2)$
- ▶ **Perfect but expensive:** Exact values but need $|S|$ memory slots



Tabular: Each state has its own value


The Problem

What if we have 1 million states? We need 1 million separate values!

Function Approximation: The Core Idea

Key Innovation

Instead of storing individual values, use a **function** with **parameters**



Why Function Approximation Works

The Magic of Shared Parameters

- ▶ **Shared parameters:** Same w computes values for ALL states

Why Function Approximation Works

The Magic of Shared Parameters

- ▶ **Shared parameters:** Same w computes values for ALL states
- ▶ **Generalization:** Similar states get similar values automatically

Why Function Approximation Works

The Magic of Shared Parameters

- ▶ **Shared parameters:** Same w computes values for ALL states
- ▶ **Generalization:** Similar states get similar values automatically
- ▶ **Efficiency:** Store only d parameters instead of $|S|$ values

Why Function Approximation Works

The Magic of Shared Parameters

- ▶ **Shared parameters:** Same w computes values for ALL states
- ▶ **Generalization:** Similar states get similar values automatically
- ▶ **Efficiency:** Store only d parameters instead of $|S|$ values

Key Benefits

- ▶ **Memory savings:** $d \ll |S|$ means huge space reduction

Why Function Approximation Works

The Magic of Shared Parameters

- ▶ **Shared parameters:** Same w computes values for ALL states
- ▶ **Generalization:** Similar states get similar values automatically
- ▶ **Efficiency:** Store only d parameters instead of $|S|$ values

Key Benefits

- ▶ **Memory savings:** $d \ll |S|$ means huge space reduction
- ▶ **Learning transfer:** Knowledge about one state helps with similar states

Why Function Approximation Works

The Magic of Shared Parameters

- ▶ **Shared parameters:** Same w computes values for ALL states
- ▶ **Generalization:** Similar states get similar values automatically
- ▶ **Efficiency:** Store only d parameters instead of $|S|$ values

Key Benefits

- ▶ **Memory savings:** $d \ll |S|$ means huge space reduction
- ▶ **Learning transfer:** Knowledge about one state helps with similar states
- ▶ **Handles continuous spaces:** Can work with infinite state spaces

Why Function Approximation Works

The Magic of Shared Parameters

- ▶ **Shared parameters:** Same w computes values for ALL states
- ▶ **Generalization:** Similar states get similar values automatically
- ▶ **Efficiency:** Store only d parameters instead of $|S|$ values

Key Benefits

- ▶ **Memory savings:** $d \ll |S|$ means huge space reduction
- ▶ **Learning transfer:** Knowledge about one state helps with similar states
- ▶ **Handles continuous spaces:** Can work with infinite state spaces

Trade-off

We gain efficiency and generalization, but lose the guarantee of exact values

Concrete Example: GridWorld

Scenario: 4×4 GridWorld (16 states)

Tabular: Need 16 separate values

Function Approximation: Maybe just 3 parameters!

(0,3)	(1,3)	(2,3)	(3,3)
(0,2)	(1,2)	(2,2)	(3,2)
(0,1)	(1,1)	(2,1)	(3,1)
(0,0)	(1,0)	(2,0)	(3,0)

GridWorld States

Function:

$$\hat{v}(x, y, w) = w_1 \cdot x + w_2 \cdot y + w_3$$

$$w = \begin{bmatrix} 0.5 \\ -0.3 \\ 2.0 \end{bmatrix}$$

$$\hat{v}(0, 0) = 0.5(0) + (-0.3)(0) + 2.0 = 2.0$$

$$\hat{v}(1, 0) = 0.5(1) + (-0.3)(0) + 2.0 = 2.5$$

$$\hat{v}(3, 3) = 0.5(3) + (-0.3)(3) + 2.0 = 2.6$$

Concrete Example: GridWorld

(0,3)	(1,3)	(2,3)	(3,3)
(0,2)	(1,2)	(2,2)	(3,2)
(0,1)	(1,1)	(2,1)	(3,1)
(0,0)	(1,0)	(2,0)	(3,0)

GridWorld States

Function:

$$\hat{v}(x, y, w) = w_1 \cdot x + w_2 \cdot y + w_3$$

$$w = \begin{bmatrix} 0.5 \\ -0.3 \\ 2.0 \end{bmatrix}$$

$$\hat{v}(0, 0) = 0.5(0) + (-0.3)(0) + 2.0 = 2.0$$

$$\hat{v}(1, 0) = 0.5(1) + (-0.3)(0) + 2.0 = 2.5$$

$$\hat{v}(3, 3) = 0.5(3) + (-0.3)(3) + 2.0 = 2.6$$

Concrete Example: GridWorld

(0,3)	(1,3)	(2,3)	(3,3)
(0,2)	(1,2)	(2,2)	(3,2)
(0,1)	(1,1)	(2,1)	(3,1)
(0,0)	(1,0)	(2,0)	(3,0)

GridWorld States

Function:

$$\hat{v}(x, y, w) = w_1 \cdot x + w_2 \cdot y + w_3$$

$$w = \begin{bmatrix} 0.5 \\ -0.3 \\ 2.0 \end{bmatrix}$$

$$\hat{v}(0, 0) = 0.5(0) + (-0.3)(0) + 2.0 = 2.0$$

$$\hat{v}(1, 0) = 0.5(1) + (-0.3)(0) + 2.0 = 2.5$$

$$\hat{v}(3, 3) = 0.5(3) + (-0.3)(3) + 2.0 = 2.6$$

Key Insight

Changing w_1 affects the value of **all** states. This is both powerful and challenging.

Update Targets in Approximation

General Update Form

$$\text{NewEstimate} \leftarrow \text{OldEstimate} + \text{StepSize}[\text{Target} - \text{OldEstimate}]$$

Update Targets in Approximation

General Update Form

$$\text{NewEstimate} \leftarrow \text{OldEstimate} + \text{StepSize}[\text{Target} - \text{OldEstimate}]$$

Update Notation: $s \mapsto u$

- ▶ s : The state being updated

Update Targets in Approximation

General Update Form

$$\text{NewEstimate} \leftarrow \text{OldEstimate} + \text{StepSize}[\text{Target} - \text{OldEstimate}]$$

Update Notation: $s \mapsto u$

- ▶ s : The state being updated
- ▶ u : The **target value** for that state

Update Targets in Approximation

General Update Form

$$\text{NewEstimate} \leftarrow \text{OldEstimate} + \text{StepSize}[\text{Target} - \text{OldEstimate}]$$

Update Notation: $s \mapsto u$

- ▶ s : The state being updated
- ▶ u : The **target value** for that state

Method-Specific Targets

- ▶ **Monte Carlo**: $s_t \mapsto G_t$ (return)

Update Targets in Approximation

General Update Form

$$\text{NewEstimate} \leftarrow \text{OldEstimate} + \text{StepSize}[\text{Target} - \text{OldEstimate}]$$

Update Notation: $s \mapsto u$

- ▶ s : The state being updated
- ▶ u : The **target value** for that state

Method-Specific Targets

- ▶ **Monte Carlo**: $s_t \mapsto G_t$ (return)
- ▶ **TD(0)**: $s_t \mapsto r_{t+1} + \gamma \hat{v}(s_{t+1}, w)$

Update Targets in Approximation

General Update Form

$$\text{NewEstimate} \leftarrow \text{OldEstimate} + \text{StepSize}[\text{Target} - \text{OldEstimate}]$$

Update Notation: $s \mapsto u$

- ▶ s : The state being updated
- ▶ u : The **target value** for that state

Method-Specific Targets

- ▶ **Monte Carlo**: $s_t \mapsto G_t$ (return)
- ▶ **TD(0)**: $s_t \mapsto r_{t+1} + \gamma \hat{v}(s_{t+1}, w)$
- ▶ **n-step TD**: $s_t \mapsto G_{t:t+n}$


Update Targets in Approximation

General Update Form

$$\text{NewEstimate} \leftarrow \text{OldEstimate} + \text{StepSize}[\text{Target} - \text{OldEstimate}]$$

Update Notation: $s \mapsto u$

- ▶ s : The state being updated
- ▶ u : The **target value** for that state




Method-Specific Targets

- ▶ **Monte Carlo**: $s_t \mapsto G_t$ (return)
- ▶ **TD(0)**: $s_t \mapsto r_{t+1} + \gamma \hat{v}(s_{t+1}, w)$
- ▶ **n-step TD**: $s_t \mapsto G_{t:t+n}$

The Generalization Challenge

Key Problem


Updating one state affects the approximation for **all other states!**



The Generalization Challenge

Key Problem

Updating one state affects the approximation for **all other states!**




Consequence

- ▶ Making one state's estimate **more accurate**

The Generalization Challenge

Key Problem

Updating one state affects the approximation for **all other states!**




Consequence

- ▶ Making one state's estimate **more accurate**
- ▶ Invariably makes others' **less accurate**

The Generalization Challenge

Key Problem

Updating one state affects the approximation for **all other states!**




Consequence

- ▶ Making one state's estimate **more accurate**
- ▶ Invariably makes others' **less accurate**
- ▶ Need to prioritize which states matter most

Defining State Importance

State Distribution $\mu(s)$


- ▶ $\mu(s) \geq 0$ for all states s



Defining State Importance

State Distribution $\mu(s)$


- ▶ $\mu(s) \geq 0$ for all states s
- ▶ $\sum_s \mu(s) = 1$ (probability distribution)



Defining State Importance

State Distribution $\mu(s)$

- ▶ $\mu(s) \geq 0$ for all states s
- ▶ $\sum_s \mu(s) = 1$ (probability distribution)
- ▶ **Higher $\mu(s)$ = more important state**



Defining State Importance

State Distribution $\mu(s)$

- ▶ $\mu(s) \geq 0$ for all states s
- ▶ $\sum_s \mu(s) = 1$ (probability distribution)
- ▶ **Higher $\mu(s)$** = **more important** state

Defining State Importance

State Distribution $\mu(s)$

- ▶ $\mu(s) \geq 0$ for all states s
- ▶ $\sum_s \mu(s) = 1$ (probability distribution)
- ▶ **Higher $\mu(s)$** = **more important** state

Common Choice: On-Policy Distribution

- ▶ $\mu(s) =$ fraction of time spent in state s under policy π

Defining State Importance

State Distribution $\mu(s)$

- ▶ $\mu(s) \geq 0$ for all states s
- ▶ $\sum_s \mu(s) = 1$ (probability distribution)
- ▶ **Higher $\mu(s)$** = **more important** state

Common Choice: On-Policy Distribution

- ▶ $\mu(s) =$ fraction of time spent in state s under policy π
- ▶ Focus learning on **frequently visited** states

The Objective Function

Mean Squared Value Error (MSVE)

$$\overline{VE}(w) = \sum_{s \in S} \mu(s) [v_\pi(s) - \hat{v}(s, w)]^2 \quad (1)$$

The Objective Function

Mean Squared Value Error (MSVE)

$$\overline{VE}(w) = \sum_{s \in S} \mu(s) [v_\pi(s) - \hat{v}(s, w)]^2 \quad (1)$$

Components

- ▶ $v_\pi(s)$: **True value** under policy π

The Objective Function

Mean Squared Value Error (MSVE)

$$\overline{VE}(w) = \sum_{s \in S} \mu(s) [v_\pi(s) - \hat{v}(s, w)]^2 \quad (1)$$

Components

- ▶ $v_\pi(s)$: **True value** under policy π
- ▶ $\hat{v}(s, w)$: **Approximate value** with weights w

The Objective Function

Mean Squared Value Error (MSVE)

$$\overline{VE}(w) = \sum_{s \in S} \mu(s) [v_\pi(s) - \hat{v}(s, w)]^2 \quad (1)$$

Components

- ▶ $v_\pi(s)$: **True value** under policy π
- ▶ $\hat{v}(s, w)$: **Approximate value** with weights w
- ▶ $\mu(s)$: **State importance** weighting

The Objective Function

Mean Squared Value Error (MSVE)

$$\overline{VE}(w) = \sum_{s \in S} \mu(s) [v_\pi(s) - \hat{v}(s, w)]^2 \quad (2)$$

The Objective Function

Mean Squared Value Error (MSVE)

$$\overline{VE}(w) = \sum_{s \in S} \mu(s) [v_\pi(s) - \hat{v}(s, w)]^2 \quad (2)$$

Important Limitation

The Objective Function

Mean Squared Value Error (MSVE)

$$\overline{VE}(w) = \sum_{s \in S} \mu(s) [v_\pi(s) - \hat{v}(s, w)]^2 \quad (2)$$

Important Limitation

- ▶ No guarantee of **global optimum**

The Objective Function

Mean Squared Value Error (MSVE)

$$\overline{VE}(w) = \sum_{s \in S} \mu(s) [v_\pi(s) - \hat{v}(s, w)]^2 \quad (2)$$

Important Limitation

- ▶ No guarantee of **global optimum**
- ▶ Usually converges to **local optimum**

The Objective Function

Mean Squared Value Error (MSVE)

$$\overline{VE}(w) = \sum_{s \in S} \mu(s) [v_\pi(s) - \hat{v}(s, w)]^2 \quad (2)$$


Important Limitation

- ▶ No guarantee of **global optimum**
- ▶ Usually converges to **local optimum**
- ▶ Quality depends on function approximation choice

Gradient Descent Intuition

What is a Gradient?


- ▶ **Slope** of a function at a specific point



Gradient Descent Intuition

What is a Gradient?


- ▶ Slope of a function at a specific point
- ▶ Points in direction of steepest increase



Gradient Descent Intuition

What is a Gradient?


- ▶ **Slope** of a function at a specific point
- ▶ Points in direction of **steepest increase**
- ▶ For function $f(w)$: $\nabla f(w) = \left(\frac{\partial f}{\partial w_1}, \dots, \frac{\partial f}{\partial w_d} \right)^T$



Gradient Descent Intuition

What is a Gradient?

- ▶ **Slope** of a function at a specific point
- ▶ Points in direction of **steepest increase**
- ▶ For function $f(w)$: $\nabla f(w) = \left(\frac{\partial f}{\partial w_1}, \dots, \frac{\partial f}{\partial w_d} \right)^T$



Gradient Descent Update

Gradient Descent Intuition

What is a Gradient?

- ▶ **Slope** of a function at a specific point
- ▶ Points in direction of **steepest increase**
- ▶ For function $f(w)$: $\nabla f(w) = \left(\frac{\partial f}{\partial w_1}, \dots, \frac{\partial f}{\partial w_d} \right)^T$

Gradient Descent Intuition

What is a Gradient?

- ▶ **Slope** of a function at a specific point
- ▶ Points in direction of **steepest increase**
- ▶ For function $f(w)$: $\nabla f(w) = \left(\frac{\partial f}{\partial w_1}, \dots, \frac{\partial f}{\partial w_d} \right)^T$

Gradient Descent Update

$$w_{t+1} = w_t - \alpha \nabla f(w_t)$$

where α is the **learning rate**

Example: Gradient Calculation - Setup

Function: $f(x, y) = 0.5x^2 + y^2$

Find the gradient at point $(x = 10, y = 10)$

Step 1: Partial Derivatives

$$\frac{\partial f}{\partial x} = x$$

$$\frac{\partial f}{\partial y} = 2y$$

Example: Gradient Calculation - Vector Form

Step 2: Gradient Vector

$$\nabla f(x, y) = \begin{bmatrix} x \\ 2y \end{bmatrix}$$

Step 3: Evaluate at (10, 10)

$$\nabla f(10, 10) = \begin{bmatrix} 10 \\ 20 \end{bmatrix}$$

Activity: 2D Gradient Descent Practice

Your Turn

Execute two steps of gradient descent starting from $(x_0, y_0) = (2, 3)$ with $\alpha = 0.1$ for:

$$f(x, y) = x^2 + 2y^2 - 4x - 6y + 10$$

Steps to follow:

1. Find the gradient $\nabla f(x, y)$
2. Calculate (x_1, y_1) after first step
3. Calculate (x_2, y_2) after second step

Solution: Step 1 - Gradient Calculation

Gradient of $f(x, y) = x^2 + 2y^2 - 4x - 6y + 10$

Taking partial derivatives:

$$\frac{\partial f}{\partial x} = 2x - 4 \tag{4}$$

$$\frac{\partial f}{\partial y} = 4y - 6 \tag{5}$$

Gradient Vector

$$\nabla f(x, y) = \begin{bmatrix} 2x - 4 \\ 4y - 6 \end{bmatrix}$$

Solution: Step 2 - First Iteration

Starting Point: $(x_0, y_0) = (2, 3)$

Evaluate gradient at $(2, 3)$:

$$\nabla f(2, 3) = \begin{bmatrix} 2(2) - 4 \\ 4(3) - 6 \end{bmatrix} = \begin{bmatrix} 0 \\ 6 \end{bmatrix}$$

Update Rule: $\theta_{new} = \theta_{old} - \alpha \nabla f(\theta_{old})$

$$\begin{bmatrix} x_1 \\ y_1 \end{bmatrix} = \begin{bmatrix} 2 \\ 3 \end{bmatrix} - 0.1 \begin{bmatrix} 0 \\ 6 \end{bmatrix} = \begin{bmatrix} 2.0 \\ 2.4 \end{bmatrix}$$

Solution: Step 3 - Second Iteration

Starting Point: $(x_1, y_1) = (2.0, 2.4)$

Evaluate gradient at $(2.0, 2.4)$:

$$\nabla f(2.0, 2.4) = \begin{bmatrix} 2(2.0) - 4 \\ 4(2.4) - 6 \end{bmatrix} = \begin{bmatrix} 0 \\ 3.6 \end{bmatrix}$$

Second Update

$$\begin{bmatrix} x_2 \\ y_2 \end{bmatrix} = \begin{bmatrix} 2.0 \\ 2.4 \end{bmatrix} - 0.1 \begin{bmatrix} 0 \\ 3.6 \end{bmatrix} = \begin{bmatrix} 2.0 \\ 2.04 \end{bmatrix}$$

Why SGD in Reinforcement Learning?

The Challenge

In RL, we want to learn value functions $v_\pi(s)$ or $q_\pi(s, a)$ but:

- ▶ We can't store a table for every state
- ▶ We need **function approximation**

Why SGD in Reinforcement Learning?

The Challenge

In RL, we want to learn value functions $v_\pi(s)$ or $q_\pi(s, a)$ but:

- ▶ We can't store a table for every state
- ▶ We need **function approximation**

Function Approximation

Represent value function with parameters: $\hat{v}(s, w)$

- ▶ w = weight vector (parameters)

Why SGD in Reinforcement Learning?

The Challenge

In RL, we want to learn value functions $v_\pi(s)$ or $q_\pi(s, a)$ but:

- ▶ We can't store a table for every state
- ▶ We need **function approximation**

Function Approximation

Represent value function with parameters: $\hat{v}(s, w)$

- ▶ w = weight vector (parameters)
- ▶ Could be linear: $\hat{v}(s, w) = w^T x(s)$

Why SGD in Reinforcement Learning?

The Challenge

In RL, we want to learn value functions $v_\pi(s)$ or $q_\pi(s, a)$ but:

- ▶ We can't store a table for every state
- ▶ We need **function approximation**

Function Approximation

Represent value function with parameters: $\hat{v}(s, w)$

- ▶ w = weight vector (parameters)
- ▶ Could be linear: $\hat{v}(s, w) = w^T x(s)$
- ▶ Or non-linear: neural networks, etc.

Why SGD in Reinforcement Learning?

The Challenge

In RL, we want to learn value functions $v_\pi(s)$ or $q_\pi(s, a)$ but:

- ▶ We can't store a table for every state
- ▶ We need **function approximation**

Function Approximation

Represent value function with parameters: $\hat{v}(s, w)$

- ▶ w = weight vector (parameters)
- ▶ Could be linear: $\hat{v}(s, w) = w^T x(s)$
- ▶ Or non-linear: neural networks, etc.

Goal

Find weights w that make $\hat{v}(s, w) \approx v_\pi(s)$ for all states

From Supervised Learning to RL

Supervised Learning Setting

- ▶ Have training data: $(x_1, y_1), (x_2, y_2), \dots, (x_n, y_n)$
- ▶ Know true labels y_i
- ▶ Minimize loss: $\sum_i L(f(x_i), y_i)$

From Supervised Learning to RL

Supervised Learning Setting

- ▶ Have training data: $(x_1, y_1), (x_2, y_2), \dots, (x_n, y_n)$
- ▶ Know true labels y_i
- ▶ Minimize loss: $\sum_i L(f(x_i), y_i)$

RL Setting (if we had true values)

- ▶ Training data: $(s_1, v_\pi(s_1)), (s_2, v_\pi(s_2)), \dots$

From Supervised Learning to RL

Supervised Learning Setting

- ▶ Have training data: $(x_1, y_1), (x_2, y_2), \dots, (x_n, y_n)$
- ▶ Know true labels y_i
- ▶ Minimize loss: $\sum_i L(f(x_i), y_i)$

RL Setting (if we had true values)

- ▶ Training data: $(s_1, v_\pi(s_1)), (s_2, v_\pi(s_2)), \dots$
- ▶ Minimize: $\sum_i [v_\pi(s_i) - \hat{v}(s_i, w)]^2$

From Supervised Learning to RL

Supervised Learning Setting

- ▶ Have training data: $(x_1, y_1), (x_2, y_2), \dots, (x_n, y_n)$
- ▶ Know true labels y_i
- ▶ Minimize loss: $\sum_i L(f(x_i), y_i)$

RL Setting (if we had true values)

- ▶ Training data: $(s_1, v_\pi(s_1)), (s_2, v_\pi(s_2)), \dots$
- ▶ Minimize: $\sum_i [v_\pi(s_i) - \hat{v}(s_i, w)]^2$
- ▶ Use gradient descent to find optimal w

From Supervised Learning to RL

Supervised Learning Setting

- ▶ Have training data: $(x_1, y_1), (x_2, y_2), \dots, (x_n, y_n)$
- ▶ Know true labels y_i
- ▶ Minimize loss: $\sum_i L(f(x_i), y_i)$

RL Setting (if we had true values)

- ▶ Training data: $(s_1, v_\pi(s_1)), (s_2, v_\pi(s_2)), \dots$
- ▶ Minimize: $\sum_i [v_\pi(s_i) - \hat{v}(s_i, w)]^2$
- ▶ Use gradient descent to find optimal w

The Problem

We don't know the true values $v_\pi(s_i)$! That's what we're trying to learn!

SGD for Value Function Approximation

Assumptions

- ▶ Approximate value function $\hat{v}(s, w)$ is **differentiable**

SGD for Value Function Approximation

Assumptions

- ▶ Approximate value function $\hat{v}(s, w)$ is **differentiable**
- ▶ Observe samples $s_t \mapsto v_\pi(s_t)$ (state and true value)

Gradient of Squared Error

$$\nabla [v_\pi(s_t) - \hat{v}(s_t, w)]^2 = -2 [v_\pi(s_t) - \hat{v}(s_t, w)] \nabla \hat{v}(s_t, w)$$

SGD for Value Function Approximation

Assumptions

- ▶ Approximate value function $\hat{v}(s, w)$ is **differentiable**
- ▶ Observe samples $s_t \mapsto v_\pi(s_t)$ (state and true value)
- ▶ States appear according to distribution μ

Gradient of Squared Error

$$\nabla [v_\pi(s_t) - \hat{v}(s_t, w)]^2 = -2 [v_\pi(s_t) - \hat{v}(s_t, w)] \nabla \hat{v}(s_t, w)$$

SGD Update Rule

$$w_{t+1} = w_t + \alpha [v_\pi(s_t) - \hat{v}(s_t, w_t)] \nabla \hat{v}(s_t, w_t)$$

SGD for Value Function Approximation

Gradient of Squared Error

$$\nabla [v_\pi(s_t) - \hat{v}(s_t, w)]^2 = -2[v_\pi(s_t) - \hat{v}(s_t, w)] \nabla \hat{v}(s_t, w)$$

SGD Update Rule

$$w_{t+1} = w_t + \alpha [v_\pi(s_t) - \hat{v}(s_t, w_t)] \nabla \hat{v}(s_t, w_t)$$

SGD for Value Function Approximation

Gradient of Squared Error

$$\nabla [v_\pi(s_t) - \hat{v}(s_t, w)]^2 = -2[v_\pi(s_t) - \hat{v}(s_t, w)] \nabla \hat{v}(s_t, w)$$

SGD Update Rule

$$w_{t+1} = w_t + \alpha [v_\pi(s_t) - \hat{v}(s_t, w_t)] \nabla \hat{v}(s_t, w_t)$$

Problem: We don't know $v_\pi(s_t)$!

The SGD Update Rule Explained

Standard SGD Update

$$\mathbf{w}_{t+1} = \mathbf{w}_t + \alpha [v_{\pi}(s_t) - \hat{v}(s_t, \mathbf{w}_t)] \nabla \hat{v}(s_t, \mathbf{w}_t)$$

The SGD Update Rule Explained

Standard SGD Update

$$\mathbf{w}_{t+1} = \mathbf{w}_t + \alpha [v_\pi(s_t) - \hat{v}(s_t, \mathbf{w}_t)] \nabla \hat{v}(s_t, \mathbf{w}_t)$$

Breaking it Down

- ▶ $v_\pi(s_t) - \hat{v}(s_t, \mathbf{w}_t)$ = **prediction error**

The SGD Update Rule Explained

Standard SGD Update

$$\mathbf{w}_{t+1} = \mathbf{w}_t + \alpha [v_\pi(s_t) - \hat{v}(s_t, \mathbf{w}_t)] \nabla \hat{v}(s_t, \mathbf{w}_t)$$

Breaking it Down

- ▶ $v_\pi(s_t) - \hat{v}(s_t, \mathbf{w}_t)$ = **prediction error**
- ▶ $\nabla \hat{v}(s_t, \mathbf{w}_t)$ = **gradient** (direction of steepest increase)

The SGD Update Rule Explained

Standard SGD Update

$$\mathbf{w}_{t+1} = \mathbf{w}_t + \alpha [v_\pi(s_t) - \hat{v}(s_t, \mathbf{w}_t)] \nabla \hat{v}(s_t, \mathbf{w}_t)$$

Breaking it Down

- ▶ $v_\pi(s_t) - \hat{v}(s_t, \mathbf{w}_t)$ = prediction error
- ▶ $\nabla \hat{v}(s_t, \mathbf{w}_t)$ = gradient (direction of steepest increase)
- ▶ α = learning rate (step size)

The SGD Update Rule Explained

Standard SGD Update

$$\mathbf{w}_{t+1} = \mathbf{w}_t + \alpha [v_\pi(s_t) - \hat{v}(s_t, \mathbf{w}_t)] \nabla \hat{v}(s_t, \mathbf{w}_t)$$

Breaking it Down

- ▶ $v_\pi(s_t) - \hat{v}(s_t, \mathbf{w}_t)$ = prediction error
- ▶ $\nabla \hat{v}(s_t, \mathbf{w}_t)$ = gradient (direction of steepest increase)
- ▶ α = learning rate (step size)

Intuition

- ▶ If $\hat{v}(s_t, \mathbf{w}_t) < v_\pi(s_t)$: error is positive ⇒ increase \mathbf{w} in gradient direction

The SGD Update Rule Explained

Standard SGD Update

$$\mathbf{w}_{t+1} = \mathbf{w}_t + \alpha [v_\pi(s_t) - \hat{v}(s_t, \mathbf{w}_t)] \nabla \hat{v}(s_t, \mathbf{w}_t)$$

Breaking it Down

- ▶ $v_\pi(s_t) - \hat{v}(s_t, \mathbf{w}_t)$ = prediction error
- ▶ $\nabla \hat{v}(s_t, \mathbf{w}_t)$ = gradient (direction of steepest increase)
- ▶ α = learning rate (step size)

Intuition

- ▶ If $\hat{v}(s_t, \mathbf{w}_t) < v_\pi(s_t)$: error is positive \Rightarrow increase \mathbf{w} in gradient direction
- ▶ If $\hat{v}(s_t, \mathbf{w}_t) > v_\pi(s_t)$: error is negative \Rightarrow decrease \mathbf{w} in gradient direction

Using Target Approximations

Solution: Unbiased Target Estimates

Replace true value $v_\pi(s_t)$ with unbiased estimate U_t

Using Target Approximations

Solution: Unbiased Target Estimates

Replace true value $v_\pi(s_t)$ with unbiased estimate U_t

General SGD Update

$$\mathbf{w}_{t+1} = \mathbf{w}_t + \alpha [U_t - \hat{v}(s_t, \mathbf{w}_t)] \nabla \hat{v}(s_t, \mathbf{w}_t)$$

Using Target Approximations

Solution: Unbiased Target Estimates

Replace true value $v_\pi(s_t)$ with unbiased estimate U_t

General SGD Update

$$\mathbf{w}_{t+1} = \mathbf{w}_t + \alpha [U_t - \hat{v}(s_t, \mathbf{w}_t)] \nabla \hat{v}(s_t, \mathbf{w}_t)$$

Target Choices

- ▶ Monte Carlo: $U_t = G_t$ (sample return)

Using Target Approximations

Solution: Unbiased Target Estimates

Replace true value $v_\pi(s_t)$ with unbiased estimate U_t

General SGD Update

$$\mathbf{w}_{t+1} = \mathbf{w}_t + \alpha [U_t - \hat{v}(s_t, \mathbf{w}_t)] \nabla \hat{v}(s_t, \mathbf{w}_t)$$

Target Choices

- ▶ Monte Carlo: $U_t = G_t$ (sample return)
- ▶ TD(0): $U_t = r_{t+1} + \gamma \hat{v}(s_{t+1}, \mathbf{w}_t)$

Using Target Approximations

Solution: Unbiased Target Estimates

Replace true value $v_\pi(s_t)$ with unbiased estimate U_t

General SGD Update

$$\mathbf{w}_{t+1} = \mathbf{w}_t + \alpha [U_t - \hat{v}(s_t, \mathbf{w}_t)] \nabla \hat{v}(s_t, \mathbf{w}_t)$$

Target Choices

- ▶ **Monte Carlo:** $U_t = G_t$ (sample return)
- ▶ **TD(0):** $U_t = r_{t+1} + \gamma \hat{v}(s_{t+1}, \mathbf{w}_t)$

Key Insight

The target U_t must be an **unbiased estimate** of $v_\pi(s_t)$ for convergence guarantees

Target Approximations: Detailed Explanation

Why Different Targets?

Each target U_t represents a different way to estimate $v_{\pi}(s_t)$:

Target Approximations: Detailed Explanation

Why Different Targets?

Each target U_t represents a different way to estimate $v_{\pi}(s_t)$:

Monte Carlo Target: $U_t = G_t$

- ▶ $G_t = r_{t+1} + \gamma r_{t+2} + \gamma^2 r_{t+3} + \dots$ (actual return)

Target Approximations: Detailed Explanation

Why Different Targets?

Each target U_t represents a different way to estimate $v_\pi(s_t)$:

Monte Carlo Target: $U_t = G_t$

- ▶ $G_t = r_{t+1} + \gamma r_{t+2} + \gamma^2 r_{t+3} + \dots$ (actual return)
- ▶ **Pro:** Unbiased estimate of $v_\pi(s_t)$

Target Approximations: Detailed Explanation

Why Different Targets?

Each target U_t represents a different way to estimate $v_\pi(s_t)$:

Monte Carlo Target: $U_t = G_t$

- ▶ $G_t = r_{t+1} + \gamma r_{t+2} + \gamma^2 r_{t+3} + \dots$ (actual return)
- ▶ **Pro:** Unbiased estimate of $v_\pi(s_t)$
- ▶ **Con:** High variance, need complete episodes

Target Approximations: Detailed Explanation

Why Different Targets?

Each target U_t represents a different way to estimate $v_\pi(s_t)$:

Monte Carlo Target: $U_t = G_t$

- ▶ $G_t = r_{t+1} + \gamma r_{t+2} + \gamma^2 r_{t+3} + \dots$ (actual return)
- ▶ **Pro:** Unbiased estimate of $v_\pi(s_t)$
- ▶ **Con:** High variance, need complete episodes

TD(0) Target: $U_t = r_{t+1} + \gamma \hat{v}(s_{t+1}, w_t)$

- ▶ Bootstrap from current estimate of next state

Target Approximations: Detailed Explanation

Why Different Targets?

Each target U_t represents a different way to estimate $v_\pi(s_t)$:

Monte Carlo Target: $U_t = G_t$

- ▶ $G_t = r_{t+1} + \gamma r_{t+2} + \gamma^2 r_{t+3} + \dots$ (actual return)
- ▶ **Pro:** Unbiased estimate of $v_\pi(s_t)$
- ▶ **Con:** High variance, need complete episodes

TD(0) Target: $U_t = r_{t+1} + \gamma \hat{v}(s_{t+1}, w_t)$

- ▶ Bootstrap from current estimate of next state
- ▶ **Pro:** Lower variance, online updates

Target Approximations: Detailed Explanation

Why Different Targets?

Each target U_t represents a different way to estimate $v_\pi(s_t)$:

Monte Carlo Target: $U_t = G_t$

- ▶ $G_t = r_{t+1} + \gamma r_{t+2} + \gamma^2 r_{t+3} + \dots$ (actual return)
- ▶ **Pro:** Unbiased estimate of $v_\pi(s_t)$
- ▶ **Con:** High variance, need complete episodes

TD(0) Target: $U_t = r_{t+1} + \gamma \hat{v}(s_{t+1}, w_t)$

- ▶ Bootstrap from current estimate of next state
- ▶ **Pro:** Lower variance, online updates
- ▶ **Con:** Biased (depends on current approximation)

Convergence Properties

When Does SGD Converge?

For convergence, we need:

1. Target U_t is an **unbiased estimate** of $v_{\pi}(s_t)$

Convergence Properties

When Does SGD Converge?

For convergence, we need:

1. Target U_t is an **unbiased estimate** of $v_{\pi}(s_t)$
2. Learning rate conditions: $\sum_t \alpha_t = \infty$ and $\sum_t \alpha_t^2 < \infty$

Convergence Properties

When Does SGD Converge?

For convergence, we need:

1. Target U_t is an **unbiased estimate** of $v_{\pi}(s_t)$
2. Learning rate conditions: $\sum_t \alpha_t = \infty$ and $\sum_t \alpha_t^2 < \infty$
3. Function approximator is **differentiable**

Convergence Properties

When Does SGD Converge?

For convergence, we need:

1. Target U_t is an **unbiased estimate** of $v_\pi(s_t)$
2. Learning rate conditions: $\sum_t \alpha_t = \infty$ and $\sum_t \alpha_t^2 < \infty$
3. Function approximator is **differentiable**

Convergence Guarantees

- **Monte Carlo:** Converges to global optimum (unbiased target)

Convergence Properties

When Does SGD Converge?

For convergence, we need:

1. Target U_t is an **unbiased estimate** of $v_\pi(s_t)$
2. Learning rate conditions: $\sum_t \alpha_t = \infty$ and $\sum_t \alpha_t^2 < \infty$
3. Function approximator is **differentiable**

Convergence Guarantees

- ▶ **Monte Carlo**: Converges to global optimum (unbiased target)
- ▶ **TD methods**: May not converge (biased target)

Convergence Properties

When Does SGD Converge?

For convergence, we need:

1. Target U_t is an **unbiased estimate** of $v_\pi(s_t)$
2. Learning rate conditions: $\sum_t \alpha_t = \infty$ and $\sum_t \alpha_t^2 < \infty$
3. Function approximator is **differentiable**

Convergence Guarantees

- ▶ **Monte Carlo**: Converges to global optimum (unbiased target)
- ▶ **TD methods**: May not converge (biased target)
- ▶ **Practice**: TD methods often work well despite theory

Convergence Properties

When Does SGD Converge?

For convergence, we need:

1. Target U_t is an **unbiased estimate** of $v_\pi(s_t)$
2. Learning rate conditions: $\sum_t \alpha_t = \infty$ and $\sum_t \alpha_t^2 < \infty$
3. Function approximator is **differentiable**

Convergence Guarantees

- ▶ **Monte Carlo**: Converges to global optimum (unbiased target)
- ▶ **TD methods**: May not converge (biased target)
- ▶ **Practice**: TD methods often work well despite theory

Key Takeaway

There's a bias-variance tradeoff in choosing targets for SGD

Gradient Monte Carlo Method

Algorithm Gradient Monte Carlo for Value Function Approximation

- 1: **Input:** Policy π , number of episodes N , step size $\alpha \in (0, 1]$
- 2: **Initialize:** $w \in \mathbb{R}^d$ arbitrarily (e.g., $w = 0$)
- 3: **for** N episodes **do**
- 4: Generate episode using π : $S_0, A_0, R_1, S_1, \dots, S_{T-1}, A_{T-1}, R_T$
- 5: **for** each step $t = 0, 1, \dots, T - 1$ **do**
- 6: $w \leftarrow w + \alpha[G_t - \hat{v}(S_t, w)]\nabla\hat{v}(S_t, w)$
- 7: **end for**
- 8: **end for**

Gradient Monte Carlo Method

Algorithm Gradient Monte Carlo for Value Function Approximation

- 1: **Input:** Policy π , number of episodes N , step size $\alpha \in (0, 1]$
- 2: **Initialize:** $w \in \mathbb{R}^d$ arbitrarily (e.g., $w = 0$)
- 3: **for** N episodes **do**
- 4: Generate episode using π : $S_0, A_0, R_1, S_1, \dots, S_{T-1}, A_{T-1}, R_T$
- 5: **for** each step $t = 0, 1, \dots, T - 1$ **do**
- 6: $w \leftarrow w + \alpha[G_t - \hat{v}(S_t, w)]\nabla\hat{v}(S_t, w)$
- 7: **end for**
- 8: **end for**

Key Features

- ▶ Uses **actual returns** G_t as targets

Gradient Monte Carlo Method

Algorithm Gradient Monte Carlo for Value Function Approximation

- 1: **Input:** Policy π , number of episodes N , step size $\alpha \in (0, 1]$
- 2: **Initialize:** $w \in \mathbb{R}^d$ arbitrarily (e.g., $w = 0$)
- 3: **for** N episodes **do**
- 4: Generate episode using π : $S_0, A_0, R_1, S_1, \dots, S_{T-1}, A_{T-1}, R_T$
- 5: **for** each step $t = 0, 1, \dots, T - 1$ **do**
- 6: $w \leftarrow w + \alpha[G_t - \hat{v}(S_t, w)]\nabla\hat{v}(S_t, w)$
- 7: **end for**
- 8: **end for**

Key Features

- ▶ Uses **actual returns** G_t as targets
- ▶ **Unbiased** - guaranteed convergence to local optimum

Gradient Monte Carlo Method

Algorithm Gradient Monte Carlo for Value Function Approximation

- 1: **Input:** Policy π , number of episodes N , step size $\alpha \in (0, 1]$
- 2: **Initialize:** $w \in \mathbb{R}^d$ arbitrarily (e.g., $w = 0$)
- 3: **for** N episodes **do**
- 4: Generate episode using π : $S_0, A_0, R_1, S_1, \dots, S_{T-1}, A_{T-1}, R_T$
- 5: **for** each step $t = 0, 1, \dots, T - 1$ **do**
- 6: $w \leftarrow w + \alpha[G_t - \hat{v}(S_t, w)]\nabla\hat{v}(S_t, w)$
- 7: **end for**
- 8: **end for**

Key Features

- ▶ Uses **actual returns** G_t as targets
- ▶ **Unbiased** - guaranteed convergence to local optimum

Semi-Gradient TD(0)

Algorithm Semi-Gradient TD(0)

- 1: **Input:** Policy π , step size $\alpha \in (0, 1]$
- 2: **Initialize:** $w \in \mathbb{R}^d$ arbitrarily
- 3: **repeat**
- 4: Initialize S
- 5: **repeat**
- 6: $A \leftarrow \pi(\cdot | S)$
- 7: Take action A , observe R, S'
- 8: $w \leftarrow w + \alpha[R + \gamma \hat{v}(S', w) - \hat{v}(S, w)]\nabla \hat{v}(S, w)$
- 9: $S \leftarrow S'$
- 10: **until** S is terminal
- 11: **until** convergence

Semi-Gradient TD(0)

Why "Semi-Gradient"?

Semi-Gradient TD(0)

Why "Semi-Gradient"?

- ▶ Target includes $\hat{v}(S', w)$ which depends on w

Semi-Gradient TD(0)

Why "Semi-Gradient"?

- ▶ Target includes $\hat{v}(S', w)$ which depends on w
- ▶ We ignore gradient w.r.t. target: **not true gradient**

Semi-Gradient TD(0)

Why "Semi-Gradient"?

- ▶ Target includes $\hat{v}(S', w)$ which depends on w
- ▶ We ignore gradient w.r.t. target: **not true gradient**
- ▶ Still converges but to different solution than true gradient

From State Values to Action Values

Extending to Control

- ▶ Need action-value function: $\hat{q}(s, a, w)$ instead of $\hat{v}(s, w)$

From State Values to Action Values

Extending to Control

- ▶ Need action-value function: $\hat{q}(s, a, w)$ instead of $\hat{v}(s, w)$
- ▶ Same principles apply with weight vector w

General Update for Action Values

$$w_{t+1} = w_t + \alpha [U_t - \hat{q}(S_t, A_t, w_t)] \nabla \hat{q}(S_t, A_t, w_t)$$

From State Values to Action Values

Extending to Control

- ▶ Need action-value function: $\hat{q}(s, a, w)$ instead of $\hat{v}(s, w)$
- ▶ Same principles apply with weight vector w

General Update for Action Values

$$w_{t+1} = w_t + \alpha [U_t - \hat{q}(S_t, A_t, w_t)] \nabla \hat{q}(S_t, A_t, w_t)$$

SARSA Target

For one-step SARSA: $U_t = R_{t+1} + \gamma \hat{q}(S_{t+1}, A_{t+1}, w_t)$

From State Values to Action Values

Extending to Control

- ▶ Need action-value function: $\hat{q}(s, a, w)$ instead of $\hat{v}(s, w)$
- ▶ Same principles apply with weight vector w

General Update for Action Values

$$w_{t+1} = w_t + \alpha [U_t - \hat{q}(S_t, A_t, w_t)] \nabla \hat{q}(S_t, A_t, w_t)$$

SARSA Target

For one-step SARSA: $U_t = R_{t+1} + \gamma \hat{q}(S_{t+1}, A_{t+1}, w_t)$

Policy Improvement

- ▶ Compute $\hat{q}(s, a, w)$ for all actions a

From State Values to Action Values

Extending to Control

- ▶ Need action-value function: $\hat{q}(s, a, w)$ instead of $\hat{v}(s, w)$
- ▶ Same principles apply with weight vector w

General Update for Action Values

$$w_{t+1} = w_t + \alpha [U_t - \hat{q}(S_t, A_t, w_t)] \nabla \hat{q}(S_t, A_t, w_t)$$

SARSA Target

For one-step SARSA: $U_t = R_{t+1} + \gamma \hat{q}(S_{t+1}, A_{t+1}, w_t)$

Policy Improvement

- ▶ Compute $\hat{q}(s, a, w)$ for all actions a
- ▶ Use ϵ -greedy: $a^* = \arg \max_a \hat{q}(s, a, w)$

Semi-Gradient SARSA Algorithm


```
1: Input:  $\hat{q} : \mathcal{S} \times \mathcal{A} \times \mathbb{R}^d \rightarrow \mathbb{R}$ , step size  $\alpha$ ,  $\epsilon > 0$ 
2: Initialize:  $w \in \mathbb{R}^d$  arbitrarily
3: repeat
4:   Initialize  $S$ 
5:    $A \leftarrow \epsilon\text{-greedy}(\hat{q}(S, \cdot, w))$ 
6:   repeat
7:     Take action  $A$ , observe  $R, S'$ 
8:     if  $S'$  is terminal then
9:        $w \leftarrow w + \alpha[R - \hat{q}(S, A, w)]\nabla\hat{q}(S, A, w)$ 
10:      break
11:    end if
12:     $A' \leftarrow \epsilon\text{-greedy}(\hat{q}(S', \cdot, w))$ 
13:     $w \leftarrow w + \alpha[R + \gamma\hat{q}(S', A', w) - \hat{q}(S, A, w)]\nabla\hat{q}(S, A, w)$ 
14:     $S \leftarrow S', A \leftarrow A'$ 
```

State Aggregation

Basic Idea

- ▶ Group **similar states** together

State Groups




State Aggregation

Basic Idea

- ▶ Group **similar states** together
- ▶ Assign **same value** to states in same group

State Groups




State Aggregation

Basic Idea

- ▶ Group **similar states** together
- ▶ Assign **same value** to states in same group
- ▶ Simplest form of function approximation

State Groups



State Aggregation

Basic Idea

- ▶ Group **similar states** together
- ▶ Assign **same value** to states in same group
- ▶ Simplest form of function approximation

Limitations

- ▶ Works only for **simple problems**

State Aggregation

Basic Idea

- ▶ Group **similar states** together
- ▶ Assign **same value** to states in same group
- ▶ Simplest form of function approximation

Limitations

- ▶ Works only for **simple problems**
- ▶ Hard to define good groupings

State Aggregation

Basic Idea

- ▶ Group **similar states** together
- ▶ Assign **same value** to states in same group
- ▶ Simplest form of function approximation

Limitations

- ▶ Works only for **simple problems**
- ▶ Hard to define good groupings
- ▶ No generalization within groups

Linear Methods

Linear Approximation

$$\hat{v}(s, w) = w^T x(s) = \sum_{i=1}^d w_i x_i(s)$$

where $x(s) = (x_1(s), x_2(s), \dots, x_d(s))^T$ is the **feature vector**

Linear Methods

Linear Approximation

$$\hat{v}(s, w) = w^T x(s) = \sum_{i=1}^d w_i x_i(s)$$

where $x(s) = (x_1(s), x_2(s), \dots, x_d(s))^T$ is the **feature vector**

Key Advantage: Simple Gradient

$$\nabla \hat{v}(s, w) = x(s)$$

Linear Methods

Linear Approximation

$$\hat{v}(s, w) = w^T x(s) = \sum_{i=1}^d w_i x_i(s)$$

where $x(s) = (x_1(s), x_2(s), \dots, x_d(s))^T$ is the **feature vector**

Key Advantage: Simple Gradient

$$\nabla \hat{v}(s, w) = x(s)$$

Linear SGD Update

$$w_{t+1} = w_t + \alpha [U_t - \hat{v}(S_t, w_t)] x(S_t)$$

Linear Methods

Linear Approximation

$$\hat{v}(s, w) = w^T x(s) = \sum_{i=1}^d w_i x_i(s)$$

where $x(s) = (x_1(s), x_2(s), \dots, x_d(s))^T$ is the **feature vector**

Key Advantage: Simple Gradient

$$\nabla \hat{v}(s, w) = x(s)$$

Linear SGD Update

$$w_{t+1} = w_t + \alpha [U_t - \hat{v}(S_t, w_t)] x(S_t)$$

Feature Construction Example

Problem Setup

State has two dimensions: $s = (s_1, s_2)$ where $s_1, s_2 \in \mathbb{R}$

Feature Construction Example

Problem Setup

State has two dimensions: $s = (s_1, s_2)$ where $s_1, s_2 \in \mathbb{R}$

Simple Linear Features

$$\mathbf{x}(s) = [s_1, s_2]^T$$

Feature Construction Example

Problem Setup

State has two dimensions: $s = (s_1, s_2)$ where $s_1, s_2 \in \mathbb{R}$

Simple Linear Features

$$x(s) = [s_1, s_2]^T$$

- ▶ **Problem:** Cannot represent interactions between s_1 and s_2

Feature Construction Example

Problem Setup

State has two dimensions: $s = (s_1, s_2)$ where $s_1, s_2 \in \mathbb{R}$

Simple Linear Features

$$x(s) = [s_1, s_2]^T$$

- ▶ **Problem:** Cannot represent interactions between s_1 and s_2
- ▶ **Problem:** If $s_1 = s_2 = 0$, then $\hat{v}(s, w) = 0$ always

Feature Construction Example

Problem Setup

State has two dimensions: $s = (s_1, s_2)$ where $s_1, s_2 \in \mathbb{R}$

Simple Linear Features

$$\mathbf{x}(s) = [s_1, s_2]^T$$

- ▶ **Problem:** Cannot represent interactions between s_1 and s_2
- ▶ **Problem:** If $s_1 = s_2 = 0$, then $\hat{v}(s, w) = 0$ always

Polynomial Features

$$\mathbf{x}(s) = [1, s_1, s_2, s_1 s_2]^T$$

Feature Construction Example

Problem Setup

State has two dimensions: $s = (s_1, s_2)$ where $s_1, s_2 \in \mathbb{R}$

Simple Linear Features

$$\mathbf{x}(s) = [s_1, s_2]^T$$

- ▶ **Problem:** Cannot represent interactions between s_1 and s_2
- ▶ **Problem:** If $s_1 = s_2 = 0$, then $\hat{v}(s, w) = 0$ always

Polynomial Features

$$\mathbf{x}(s) = [1, s_1, s_2, s_1 s_2]^T$$

- ▶ **Benefit:** Can represent interactions

Feature Construction Example

Problem Setup

State has two dimensions: $s = (s_1, s_2)$ where $s_1, s_2 \in \mathbb{R}$

Simple Linear Features

$$\mathbf{x}(s) = [s_1, s_2]^T$$

- ▶ **Problem:** Cannot represent interactions between s_1 and s_2
- ▶ **Problem:** If $s_1 = s_2 = 0$, then $\hat{v}(s, w) = 0$ always

Polynomial Features

$$\mathbf{x}(s) = [1, s_1, s_2, s_1 s_2]^T$$

- ▶ **Benefit:** Can represent interactions

Feature Construction Example

Problem Setup

State has two dimensions: $s = (s_1, s_2)$ where $s_1, s_2 \in \mathbb{R}$

Simple Linear Features

$$\mathbf{x}(s) = [s_1, s_2]^T$$

- ▶ **Problem:** Cannot represent interactions between s_1 and s_2
- ▶ **Problem:** If $s_1 = s_2 = 0$, then $\hat{v}(s, w) = 0$ always

Polynomial Features

$$\mathbf{x}(s) = [1, s_1, s_2, s_1 s_2]^T$$

- ▶ **Benefit:** Can represent interactions

Higher-Order Polynomial

$$\mathbf{x}(s) = [1, s_1, s_2, s_1 s_2, s_1^2, s_2^2, s_1^2 s_2, s_1 s_2^2, s_1^2 s_2^2]^T$$